

Jiarui Fang (方佳瑞)

☎ + (86) 13717819702

✉ fangjiarui123@gmail.com

🏠 China Technology Exchange Center, Haidian Beijing, China

🌐 <https://github.com/feifeibear>

📅 **Update Date:** September 13, 2021

OBJECTIVE: Research Scientist / Software Engineer in **High Performance Computing Systems**.

Work Experience

- **WeChat, Tencent** Beijing, China
Senior Software Engineer at WeChat AI July 2019 - Present
Mentor : Dr. Cheng Niu [link]
- **National Supercomputing Center in Wuxi** Wuxi, Jiangsu, China
Ph.D Research Intern at R&D Center August 2017 - May 2019
Mentor : Prof. Haohuan Fu

Education

- **Tsinghua University** Beijing, China
Ph.D. in Department of Computer Science & Technology July 2019
Advisor : Prof. Guangwen Yang, Coadvisor: Prof. Haohuan Fu
- **University of California, Davis** Davis, CA, USA
Visiting Scholar in Department of Computer Science Engineering August 2017 - August 2018
Advisor : Assistant Prof. Cho-Jui Hsieh [link]
- **Beijing University of Posts and Telecommunications** Beijing, China
B.S. in Department of Computer Science & Technology June 2014
Ranking 6th **top 2%** among 300 students (Honored 2014 Outstanding Graduate of Beijing)

Project Highlights

- **High Performance Offline and Online Systems for NLP Applications**
Wechat AI, Tencent July 2019 - Now
Development of TurboTransformers [link], which is an efficient GPU serving system for transformer models. Development of PatrickStar, which is a parallel training system for pre-trained models. Both systems surpass the SOTA performance at the time they are publicly released.
- **Development of Basic Modules In WeChat App**
Wechat AI, Tencent July 2019 - Now
Development of WeChat Input Method Engine, and WeChat Translation System, which are accessible in WeChat, an App of over 1 billion monthly active users in China.
- **Large-scale Deep Learning Training (DL) System for GPU Supercomputer**
University of California, Davis September 2017 - August 2018
Development of RedSync – a distributed data-parallel Deep Learning training system using gradient pruning and quantization. When scaled up to 128 GPUs, the RedSync brought significant performance improvements to DNNs that were previously considered hard to scale.
- **High Performance Deep Learning System for the Sunway TaihuLight**
National Supercomputing Center in Wuxi April 2016 - August 2019
Development of a deep learning framework from zero on the Sunway TaihuLight, which is based on the innovative SW26010 many-core processors and ranked **No.1 on the 47th-50th Top500 Supercomputer lists**.

1. Designed the swGEMM – a GGeneral Matrix Multiplication (GEMM) library based on SW26010. Core code handwritten by assembly, reaching 97% of peak performance. Significant speedups (2-10x) were achieved by applying swGEMM instead of default BLAS to deep learning applications.
2. Designed the swDNN – a library provides APIs for mainstream DL operator (CONV, LSTM, FC, BN and activations). Regarding the most complicated CONV ops, three parallel schemes were designed for the special SW26010 many-core architecture, i.e. explicit GEMM, implicit GEMM, and Winograd. The computing efficiency of swDNN exceeded cuDNNv7.5 running on Tesla K40.
3. Designed the swATOP – an end-to-end automated framework that optimizes complex parallel DL operator code on SW26010. By reading several lines of DSL statements, swATOP can automatically generate code that exceeds manual optimization performance.
4. Designed the swCaffe – an MPI-based deep learning framework on the Sunway TaihuLight. Synchronization employed an innovative topology-aware MPI Allreduce method which is 10x faster than the default MPI Allreduce on 1024 nodes.

• High Performance Scientific Computing Applications

Department of Earth System Science, Tsinghua University

February 2014 - March 2016

1. Proposed a generalized cache-friendly design based on NVIDIA GPUs and Intel Xeon Phis for complex spatially-variable coefficient (CSVC) stencils. Gained 4x speedup in the seismic imaging software (GeoEast-Lightning) used by China National Petroleum Corporation.
2. Accelerated a serial of scientific applications on different HPC platforms, including: transient electromagnetic simulation on CPU cluster; remote sensing data analysis with SVM on Intel Xeon Phi; Community Earth System Model (CESM) and crop modeling on Sunway TaihuLight.

First Author Publications¹

- Jiarui Fang, Yang Yu, Shenggui Li, Yang You, Jie Zhou, **PatrickStar: Parallel Training of Pre-trained Models via a Chunk-based Memory Management**, arXiv preprint. [pdf] .
- Jiarui Fang, Yang Yu, Chengduo. Zhao, Jie Zhou, **TurboTransformers: An Efficient GPU Serving System For Transformer Models**, Proceedings of the 26th ACM SIGPLAN Symposium on Principles and Practice of Parallel (PPoPP 2021), p389-402. [pdf] .
- Jiarui Fang, **Parallel Deep Learning Training System on on Sunway TaihuLight**, Ph.D Dissertation of Tsinghua University. (in Chinese)[pdf] .
- Jiarui Fang, Haohuan Fu, Guangwen Yang, Cho-Jui Hsieh, **RedSync : Reducing Synchronization Traffic for Distributed Deep Learning**. Journal of Parallel and Distributed Computing (JPDC), Volume 133, November 2019, p30-39. [arXiv][pdf].
- Wei Gao*, Jiarui Fang*, Wenlai Zhao, Jinzhe Yang, Long Wang, Lin Gan, Haohuan Fu, Guangwen Yang. **swATOP: Automatically Optimizing Deep Learning Operators on SW26010 Many-Core Processor**. Proceedings of the 48th International Conference on Parallel Processing (ICPP, 2019). p1-10. (* equal contribution) [pdf] .
- Jiarui Fang*, and Li, Liandeng* and Fu, Haohuan and Jiang, Jinlei and Zhao, Wenlai and He, Conghui and You, Xin and Yang, Guangwen. **swCaffe: a Parallel Framework for Accelerating Deep Learning Applications on Sunway TaihuLight**, EEE Cluster (Cluster, 2018), p413-422. [pdf]. (* equal contribution).
- Jiarui Fang, Haohuan Fu, Wenlai Zhao, Bingwei Chen, Weijie Zheng, and Guangwen Yang. **swDNN: A library for Accelerating Deep Learning Applications on Sunway Taihulight**. In Parallel and Distributed Processing Symposium (IPDPS, 2017), p615–624. [pdf]
- Jiarui Fang, Haohuan Fu, Guangwen Yang. **Cache-friendly Design for Complex Spatially-variable Coefficient Stencils on Many-core Architectures**. IEEE 23rd International Conference on High Performance Computing, Data, and Analytics (HiPC, 2016),p222-p231, [pdf]

¹Google Scholar: <https://scholar.google.com/citations?user=xibpfkQAAAAJ&hl=en>

- **Jiarui Fang**, Haohuan Fu, He Zhang, Wei Wu, Nanxun Dai, Lin Gan, Guangwen Yang. **Optimizing Complex Spatially-Variant Coefficient Stencils for Seismic Modeling on GPU**. IEEE 21st International Conference on Parallel and Distributed Systems (ICPADS, 2015), p641-p648 [pdf]
- **Jiarui Fang**, Haohuan Fu, Guangwen Yang, Wei Wu, Nanxun Dai. **GPU-based explicit time evolution method**. The 84th Society of Exploration Geophysicists Technical Program Expanded Abstracts (SEG), p3549-p3553, New Orleans, USA, 2015 [pdf]

Skills

- **Good at English:** GRE 322 (Verbal 153, October 2012)
- **Programing Language:** C/C++, CUDA, Python

References

- **Guangwen Yang**
Professor in Department of Computer Science, Tsinghua University,
Director of the National Supercomputing Center in Wuxi.
Email: ygw@tsinghua.edu.cn
- **Haohuan Fu**
Professor in Department of Earth Science, Tsinghua University,
Deputy Director of the National Supercomputing Center in Wuxi.
Email: haohuan@tsinghua.edu.cn
- **Cho-Jui Hsieh**
Assistant Professor in Department of Computer Science, University of California, Los Angeles.
Email: chohsieh@cs.ucla.edu